

eda_cbm

June 6, 2024

1 Exploratory data analysis

For this notebook, I borrowed code from Sadegh Bolouki's [Bank Account Fraud Detection: EDA and Model](#).

```
[ ]: # Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.colors
import seaborn as sns

# Define custom colors/cmmaps/palettes for visualization purposes.
denim='#6F8FAF'
salmon='#FA8072'
slate_gray = '#404040'
cmap=matplotlib.colors.LinearSegmentedColormap.from_list("",[denim,salmon])
palette = 'colorblind'
sns.set_style('darkgrid')

from sklearn.model_selection import train_test_split
```

```
[ ]: # Data loading
df = pd.read_csv('../data/Base.csv')

target = 'fraud_bool'

X = df.drop(target,axis=1)
y = df[target]
```

```
[ ]: # Train/test split with stratified sampling

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
↪test_size=0.2, random_state=42)

train = pd.concat([X_train, y_train], axis=1).copy()
train_copy = pd.concat([X_train, y_train], axis=1).copy()
```

```
[ ]: feature_descriptions = {
    'income' : "Annual income of the applicant (in decile form). Ranges between
    ↪ [0.1, 0.9].",
    'name_email_similarity' : "Metric of similarity between email and
    ↪ applicant's name. Higher values represent higher similarity. Ranges between
    ↪ [0, 1].",
    'prev_address_months_count' : "Number of months in previous registered
    ↪ address of the applicant, i.e. the applicant's previous residence, if
    ↪ applicable. Ranges between [-1, 380] months (-1 is a missing value).",
    'current_address_months_count' : "Months in currently registered address of
    ↪ the applicant. Ranges between [-1, 429] months (-1 is a missing value).",
    'customer_age' : "Applicant's age in years, rounded to the decade. Ranges
    ↪ between [10, 90] years.",
    'days_since_request' : "Number of days passed since application was done.
    ↪ Ranges between [0, 79] days.",
    'intended_balcon_amount' : "Initial transferred amount for application.
    ↪ Ranges between [-16, 114] (negatives are missing values).",
    'payment_type' : "Credit payment plan type. 5 possible (anonymized) values.
    ↪",
    'zip_count_4w' : "Number of applications within same zip code in last 4
    ↪ weeks. Ranges between [1, 6830].",
    'velocity_6h' : "Velocity of total applications made in last 6 hours i.e.,
    ↪ average number of applications per hour in the last 6 hours. Ranges between
    ↪ [-175, 16818].",
    'velocity_24h' : "Velocity of total applications made in last 24 hours i.e.
    ↪, average number of applications per hour in the last 24 hours. Ranges
    ↪ between [1297, 9586].",
    'velocity_4w' : "Velocity of total applications made in last 4 weeks, i.e.,
    ↪ average number of applications per hour in the last 4 weeks. Ranges between
    ↪ [2825, 7020].",
    'bank_branch_count_8w' : "Number of total applications in the selected bank
    ↪ branch in last 8 weeks. Ranges between [0, 2404].",
    'date_of_birth_distinct_emails_4w' : "Number of emails for applicants with
    ↪ same date of birth in last 4 weeks. Ranges between [0, 39].",
    'employment_status' : "Employment status of the applicant. 7 possible
    ↪ (anonymized) values.",
    'credit_risk_score' : "Internal score of application risk. Ranges between
    ↪ [-191, 389].",
    'email_is_free' : "Domain of application email (either free or paid).",
    'housing_status' : "Current residential status for applicant. 7 possible
    ↪ (anonymized) values.",
    'phone_home_valid' : "Validity of provided home phone.",
    'phone_mobile_valid' : "Validity of provided mobile phone.",
    'bank_months_count' : "How old is previous account (if held) in months.
    ↪ Ranges between [-1, 32] months (-1 is a missing value).",
}
```

```

    'has_other_cards' : "If applicant has other cards from the same banking
    ↪company.",
    'proposed_credit_limit' : "Applicant's proposed credit limit. Ranges
    ↪between [200, 2000].",
    'foreign_request' : "If origin country of request is different from bank's
    ↪country.",
    'source' : "Online source of application. Either browser (INTERNET) or app
    ↪(TELEAPP).",
    'session_length_in_minutes' : "Length of user session in banking website in
    ↪minutes. Ranges between [-1, 107] minutes (-1 is a missing value).",
    'device_os' : "Operative system of device that made request. Possible
    ↪values are: Windows, macOS, Linux, X11, or other.",
    'keep_alive_session' : "User option on session logout.",
    'device_distinct_emails_8w' : "Number of distinct emails in banking website
    ↪from the used device in last 8 weeks. Ranges between [-1, 2] emails (-1 is a
    ↪missing value).",
    'device_fraud_count' : "Number of fraudulent applications with used device.
    ↪Ranges between [0, 1].",
    'month' : "Month where the application was made. Ranges between [0, 7].",
    'fraud_bool' : "If the application is fraudulent or not."
}

```

```

[ ]: # Numerical (continuous/discrete) and categorical features

num_feats = X_train.select_dtypes(include='number').columns.tolist()

thresh = 25

cont_feats = [feat for feat in num_feats if df[feat].nunique() >= thresh]
disc_feats = [feat for feat in num_feats if df[feat].nunique() < thresh]

cat_feats = X_train.select_dtypes(exclude='number').columns.tolist()

print(f'Features: {X_train.shape[1]}\n\n\
Continuous: {len(cont_feats)}\n\
{cont_feats}\n\n\
Discrete: {len(disc_feats)}\n\
{disc_feats}\n\n\
Categorical: {len(cat_feats)}\n\
{cat_feats}')

```

Features: 31

Continuous: 14

```

['name_email_similarity', 'prev_address_months_count',
'current_address_months_count', 'days_since_request', 'intended_balcon_amount',
'zip_count_4w', 'velocity_6h', 'velocity_24h', 'velocity_4w',

```

```
'bank_branch_count_8w', 'date_of_birth_distinct_emails_4w', 'credit_risk_score',  
'bank_months_count', 'session_length_in_minutes']
```

Discrete: 12


```
['income', 'customer_age', 'email_is_free', 'phone_home_valid',  
'phone_mobile_valid', 'has_other_cards', 'proposed_credit_limit',  
'foreign_request', 'keep_alive_session', 'device_distinct_emails_8w',  
'device_fraud_count', 'month']
```

Categorical: 5

```
['payment_type', 'employment_status', 'housing_status', 'source', 'device_os']
```

```
[ ]: # Missing values  
# The datasheet details how the missing values are represented  
  
cols_missing_neg1 = ['prev_address_months_count',  
                     'current_address_months_count',  
                     'bank_months_count',  
                     'session_length_in_minutes',  
                     'device_distinct_emails_8w']  
  
X_train[cols_missing_neg1] = X_train[cols_missing_neg1].replace(-1,np.nan)  
  
col_missing_neg = 'intended_balcon_amount'  
  
X_train[col_missing_neg] = X_train[col_missing_neg]\  
    .apply(lambda x: np.nan if x < 0 else x)  
  
# Missing values by feature  
  
null_X_train = X_train.isna().sum()/len(X_train)*100  
  
fig = plt.figure(figsize=(6.4,4.8))  
ax = null_X_train.loc[null_X_train>0].sort_values()\  
    .plot(kind='bar',title='Percentage of missing values')  
  
for p in ax.patches:  
    ax.annotate(f'{p.get_height():.2f}%',  
                (p.get_x() + p.get_width() / 2., p.get_height()),  
                ha='center', va='bottom', xytext=(0,0), textcoords='offset_↵  
↵points')  
  
ax.set_ylabel('missing %')  
ax.set_xlabel('feature')  
  
ax.xaxis.grid(False)
```

```
plt.show()
```



```
[ ]: # Customized description and plots for any given feature
```

```
def summary(feats):  
    if feat in cont_feats:  
        cont_summary(feats)  
        cont_plots(feats)
```

```

elif feat in disc_feats:
    disc_summary(feat)
    disc_plots(feat)
else:
    cat_summary(feat)
    cat_plots(feat)

missing_flag_plot(feat)

return

# Customized description for continuous features

def cont_summary(feat):

    # Create an empty summary
    columns = ['dtype', 'count', 'unique', 'top_value_counts', 'missing_count',
               'missing_percentage', 'mean', 'std', 'min', 'median', 'max',
               'corr_with_target']
    summary = pd.DataFrame(index=[feat], columns=columns, dtype=float)

    # Pull the feature column in question
    col = X_train[feat].copy()

    # Basic statistics using the original describe method
    summary.loc[feat, ['count', 'mean', 'std', 'min', 'median', 'max']] \
        = col.describe(percentiles=[.5]).values.transpose()

    # Number of unique values
    summary.loc[feat, 'unique'] = col.nunique()

    # Missing values count
    summary.loc[feat, 'missing_count'] = col.isnull().sum()

    # Missing values percentage
    summary.loc[feat, 'missing_percentage'] = col.isnull().sum()/len(col)*100

    # Correlation with target
    summary.loc[feat, 'corr_with_target'] = col.corr(y_train)

    int_cols = ['count', 'unique', 'missing_count']
    summary[int_cols] = summary[int_cols].astype(int)
    summary = summary.round(2).astype(str)

    # Top 3 value_counts
    value_counts = X_train[feat].value_counts().head(3)
    value_counts.index = value_counts.index.astype(float).to_numpy().round(2)

```

```

summary.loc[feat, 'top_value_counts'] = str(value_counts.to_dict())

# Data type
summary.loc[feat, 'dtype'] = col.dtypes

return display(summary)

# Customized plots for continuous features

def cont_plots(feat, bins='auto'):

    n_cols = 3
    fig, axes = plt.subplots(1, n_cols, figsize=(6.4*n_cols, 4.8))

    # Histogram
    sns.histplot(data=X_train,
                  x=feat,
                  bins=bins,
                  ax=axes[0],
                  color=slate_gray)

    # Box plots with the target as hue
    sns.boxplot(data=X_train,
                 x=feat,
                 y=y_train,
                 ax=axes[1],
                 palette=palette,
                 orient='h')

    # KDE plots with the target as hue
    sns.kdeplot(data=X_train,
                 x=feat,
                 hue=y_train,
                 palette=palette,
                 fill=True,
                 common_norm=False,
                 ax=axes[2])

    axes[0].title.set_text('Histogram')
    axes[1].title.set_text('Box Plots')
    axes[2].title.set_text('KDE Plots')

    fig.tight_layout()
    plt.show()
    return

# Customized description for discrete features

```

```

def disc_summary(feat):

    # Create an empty summary
    columns = ['dtype', 'count', 'unique', 'missing_count', 'missing_percentage',
               'mean', 'std', 'min', 'median', 'max', 'cv', 'corr_with_target']
    summary = pd.DataFrame(index=[feat], columns=columns, dtype=float)

    # Pull the feature column in question
    col = X_train[feat].copy()

    # Basic statistics using the original describe method
    summary.loc[feat, ['count', 'mean', 'std', 'min', 'median', 'max']] \
    = col.describe(percentiles=[.5]).values.transpose()

    # Number of unique values
    summary.loc[feat, 'unique'] = col.nunique()

    # Coefficient of Variation (CV)
    summary.loc[feat, 'cv'] = np.NaN if not col.mean() else col.std()/col.mean()

    # Missing values count
    summary.loc[feat, 'missing_count'] = col.isnull().sum()

    # Missing values percentage
    summary.loc[feat, 'missing_percentage'] = col.isnull().sum()/len(col)*100

    # Correlation with target
    summary.loc[feat, 'corr_with_target'] = col.corr(y_train)

    int_cols = ['count', 'unique', 'missing_count']
    summary[int_cols] = summary[int_cols].astype(int)
    summary = summary.round(2).astype(str)

    # Data type
    summary.loc[feat, 'dtype'] = col.dtypes

    return display(summary)

# Customized plots for discrete features
def disc_plots(feat):

    col = X_train[feat].copy()

```



```

n_rows = 1
n_cols = 2

fig, axes = plt.subplots(n_rows, n_cols, figsize=(6.4 * n_cols, 4.8 *
↪n_rows))

# Sort unique values
unique_values = col.dropna().unique()
unique_values.sort()

# Value counts
val_counts = col.dropna().value_counts()
val_counts = val_counts.reindex(unique_values)
val_counts_pct = val_counts/len(col)*100

# Countplot
sns.countplot(x=col, order=unique_values, palette=palette, ax=axes[0])
axes[0].xaxis.grid(False)

# Show count value if rare (less than 1%)
lp_thresh = 1
for i, p in enumerate(axes[0].patches):
    pct = val_counts_pct.iloc[i]
    axes[0].annotate(f'{pct:.2f}%',
                    (p.get_x() + p.get_width() / 2., p.get_height()),
                    ha='center', va='bottom', xytext=(0,0),
                    textcoords='offset points')
    if pct < lp_thresh:
        axes[0].annotate(val_counts.iloc[i],
                        (p.get_x() + p.get_width() / 2., p.get_height()),
                        ha='center', va='bottom', xytext=(0,10),
                        textcoords='offset points',color='red')

# Barplot
df = pd.concat([X_train,y_train],axis=1).groupby(feats)[target].mean()*100
df = df.reindex(unique_values) # Reindex to match the order
sns.barplot(x=df.index, y=df.values, palette=palette, ax=axes[1])
axes[1].set_ylabel('Fraud %')
axes[1].xaxis.grid(False)

fig.tight_layout()
plt.show()

return

# Customized description for categorical features

```

```

def cat_summary(feat):

    # Create an empty summary
    columns = ['dtype', 'count', 'unique', 'missing_count', '
↳'missing_percentage']
    summary = pd.DataFrame(index=[feat], columns=columns, dtype=float)

    # Pull the feature column in question
    col = X_train[feat].copy()

    # Count
    summary.loc[feat, 'count'] = col.count()

    # Number of unique values
    summary.loc[feat, 'unique'] = col.nunique()

    # Missing values count
    summary.loc[feat, 'missing_count'] = col.isnull().sum()

    # Missing values percentage
    summary.loc[feat, 'missing_percentage'] = col.isnull().sum()/len(col)*100

    int_cols = ['count', 'unique', 'missing_count']
    summary[int_cols] = summary[int_cols].astype(int)
    summary = summary.round(2).astype(str)

    # Data type
    summary.loc[feat, 'dtype'] = col.dtypes

    return display(summary)

# Customized plots for categorical features

def cat_plots(feat):

    col = X_train[feat].copy()

    n_rows = 1
    n_cols = 2

    fig, axes = plt.subplots(n_rows, n_cols, figsize=(6.4 * n_cols, 4.8 *
↳n_rows))

    # Value counts
    val_counts = col.dropna().value_counts()

    # Unique values

```

```

unique_values = val_counts.index

# Countplot with sorted order
sns.countplot(x=col, order=unique_values, palette=palette, ax=axes[0])
axes[0].xaxis.grid(False)

val_counts_pct = val_counts/len(col)*100

# Show count value if rare (less than 1%)
lp_thresh = 1
for i, p in enumerate(axes[0].patches):
    pct = val_counts_pct.iloc[i]
    axes[0].annotate(f'{pct:.2f}%',
                    (p.get_x() + p.get_width() / 2., p.get_height()),
                    ha='center', va='bottom', xytext=(0,0),
                    textcoords='offset points')
    if pct < lp_thresh:
        axes[0].annotate(val_counts.iloc[i],
                        (p.get_x() + p.get_width() / 2., p.get_height()),
                        ha='center', va='bottom', xytext=(0,10),
                        textcoords='offset points', color='red')

# Barplot with the same order
df = pd.concat([X_train, y_train], axis=1).groupby(feats)[target].mean()*100
sns.barplot(x=df.index, y=df.values, order=unique_values, palette=palette,
↪ ax=axes[1])
axes[1].set_ylabel('Fraud %')
axes[1].xaxis.grid(False)

fig.tight_layout()
plt.show()

return

# Plot for the missing flag associated with a feature

def missing_flag_plot(feats):
    col = X_train[feats].isnull().astype(int)

    if not col.sum():
        return

    df = (pd.concat([col, y_train], axis=1).groupby(feats).mean()*100).
↪ reset_index()
    cols = [f'MISSING_FLAG_{feats}', 'Fraud %']
    df.columns = cols
    fig = plt.figure(figsize=(6.4, 4.8))

```

```
ax = sns.barplot(data=df,x=cols[0], y=cols[1], palette=palette)

fig.tight_layout()
plt.show()

return
```

1.1 Continuous features

Applying summary function to all continuous features:

```
[ ]: for feat in cont_feats:
    print(f"\033[1m Feature:\033[0m '{feat}'\n")
    print(f'\033[1m Description:\033[0m {feature_descriptions[feat]}')
    summary(feat)
    print('-'*45,'\n')
```


Feature: 'name_email_similarity'

Description: Metric of similarity between email and applicant's name.
Higher values represent higher similarity. Ranges between [0, 1].

	dtype	count	unique	top_value_counts \
name_email_similarity	float64	800000	799289	{0.23: 2, 0.26: 2, 0.78: 2}

	missing_count	missing_percentage	mean	std	min \
name_email_similarity	0		0.0	0.49	0.29 0.0

	median	max	corr_with_target
name_email_similarity	0.49	1.0	-0.04



Feature: 'prev_address_months_count'

Description: Number of months in previous registered address of the applicant, i.e. the applicant's previous residence, if applicable. Ranges between [-1, 380] months (-1 is a missing value).

```

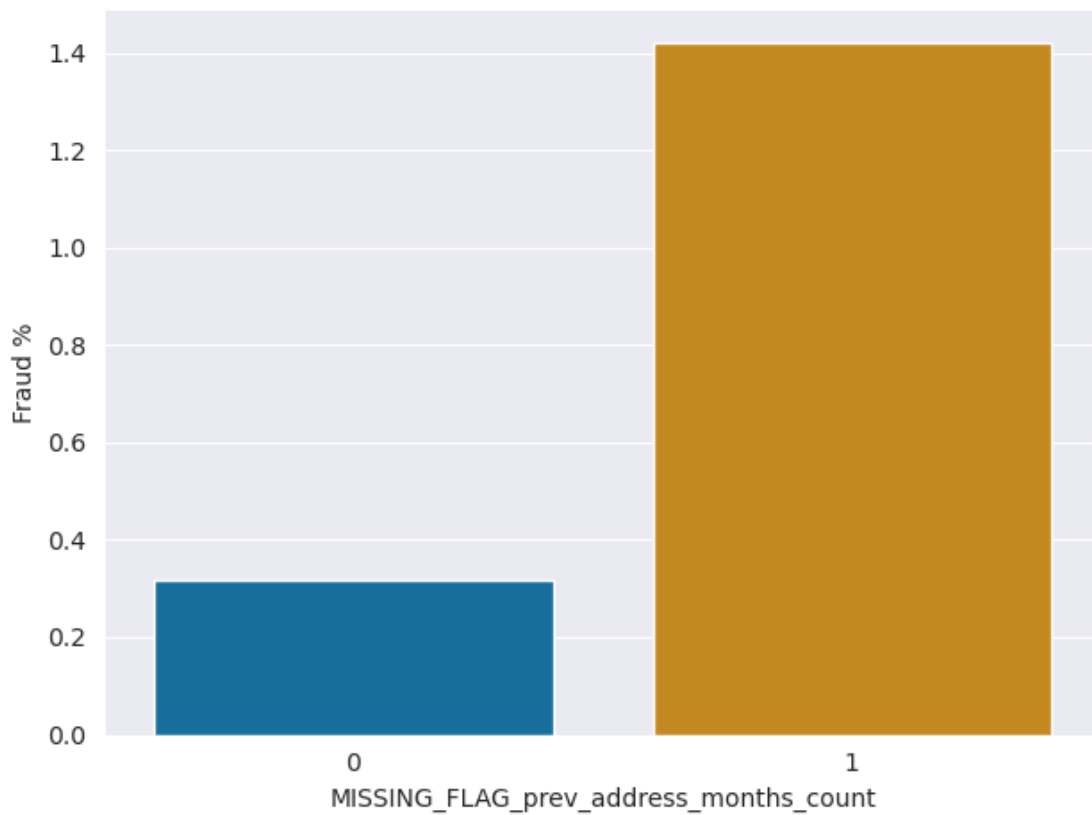

dtype    count  unique  \
prev_address_months_count  float64  229475    373

top_value_counts  missing_count  \
prev_address_months_count  {11.0: 9204, 28.0: 8246, 29.0: 8079}    570525

missing_percentage    mean    std  min  median    max  \
prev_address_months_count    71.32  60.71  63.55  5.0    34.0  383.0

corr_with_target
prev_address_months_count    0.02

```



Feature: 'current_address_months_count'


Description: Months in currently registered address of the applicant.
Ranges between [-1, 429] months (-1 is a missing value).

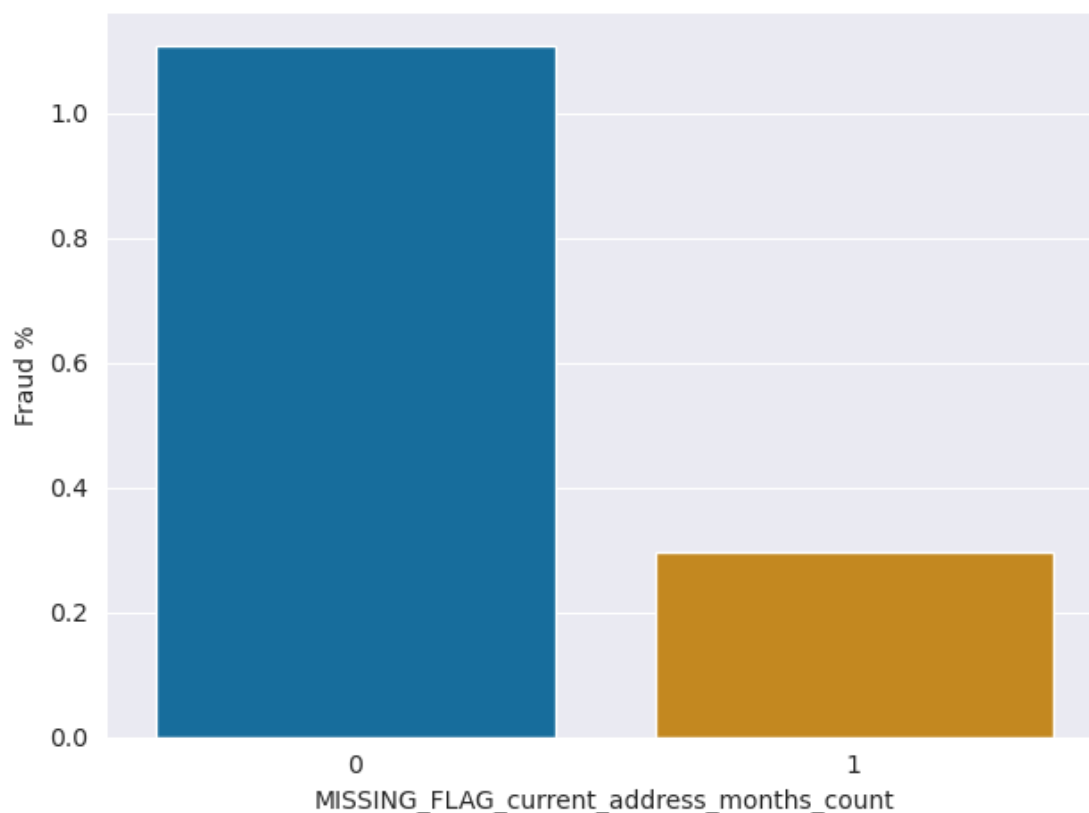
	dtype	count	unique	\
current_address_months_count	float64	796609	420	

	top_value_counts	\
current_address_months_count	{6.0: 13001, 7.0: 12952, 8.0: 12874}	

	missing_count	missing_percentage	mean	std	\
current_address_months_count	3391	0.42	86.99	88.39	

	min	median	max	corr_with_target
current_address_months_count	0.0	53.0	428.0	0.03






Feature: 'days_since_request'

Description: Number of days passed since application was done. Ranges between [0, 79] days.

	dtype	count	unique	top_value_counts	missing_count	\
days_since_request	float64	800000	793121	{0.01: 3, 0.02: 3}	0	

	missing_percentage	mean	std	min	median	max	\
days_since_request	0.0	1.02	5.38	0.0	0.02	78.46	

	corr_with_target
days_since_request	0.0



Feature: 'intended_balcon_amount'


Description: Initial transferred amount for application. Ranges between [-16, 114] (negatives are missing values).


	dtype	count	unique \
intended_balcon_amount	float64	206132	205910

	top_value_counts	missing_count \
intended_balcon_amount	{39.07: 2, 36.04: 2, 15.63: 2}	593868

	missing_percentage	mean	std	min	median	max \
intended_balcon_amount	74.23	36.57	23.22	0.0	32.45	112.96

	corr_with_target
intended_balcon_amount	0.01






Feature: 'zip_count_4w'

Description: Number of applications within same zip code in last 4 weeks. Ranges between [1, 6830].

	dtype	count	unique	top_value_counts \
zip_count_4w	int64	800000	6268	{1062.0: 645, 1020.0: 644, 941.0: 622}

	missing_count	missing_percentage	mean	std	min	median \
zip_count_4w	0	0.0	1572.77	1005.23	1.0	1263.0

	max	corr_with_target
zip_count_4w	6700.0	0.01




Feature: 'velocity_6h'

Description: Velocity of total applications made in last 6 hours i.e., average number of applications per hour in the last 6 hours. Ranges between [-175, 16818].

	dtype	count	unique	top_value_counts \
velocity_6h	float64	800000	799150	{2985.72: 2, 5900.35: 2, 8769.6: 2}

	missing_count	missing_percentage	mean	std	min \
velocity_6h	0	0.0	5664.02	3009.68	-170.6

	median	max	corr_with_target
velocity_6h	5316.3	16715.57	-0.02




Feature: 'velocity_24h'

Description: Velocity of total applications made in last 24 hours i.e., average number of applications per hour in the last 24 hours. Ranges between [1297, 9586].

	dtype	count	unique	top_value_counts \
velocity_24h	float64	800000	799310	{5082.33: 3, 4667.86: 3, 3906.42: 3}

	missing_count	missing_percentage	mean	std	min	\
velocity_24h	0	0.0	4770.23	1479.5	1320.28	

	median	max	corr_with_target
velocity_24h	4750.8	9506.9	-0.01




Feature: 'velocity_4w'

Description: Velocity of total applications made in last 4 weeks, i.e., average number of applications per hour in the last 4 weeks. Ranges between [2825, 7020].

	dtype	count	unique	top_value_counts	\
velocity_4w	float64	800000	798908	{5466.87: 3, 3137.5: 2, 4334.23: 2}	

	missing_count	missing_percentage	mean	std	min	\
velocity_4w	0	0.0	4856.0	919.62	2825.75	

	median	max	corr_with_target
velocity_4w	4913.54	6994.76	-0.01



Feature: 'bank_branch_count_8w'

Description: Number of total applications in the selected bank branch in last 8 weeks. Ranges between [0, 2404].

```


dtype    count unique \
bank_branch_count_8w  int64  800000    2318

top_value_counts missing_count \
bank_branch_count_8w  {1.0: 119074, 0.0: 115535, 2.0: 46310}          0

missing_percentage    mean    std  min median    max \
bank_branch_count_8w      0.0  184.3  459.48  0.0    9.0  2385.0

corr_with_target
bank_branch_count_8w      -0.01

```



Feature: 'date_of_birth_distinct_emails_4w'

Description: Number of emails for applicants with same date of birth in last 4 weeks. Ranges between [0, 39].

```


dtype    count unique \
date_of_birth_distinct_emails_4w  int64  800000    40

top_value_counts \
date_of_birth_distinct_emails_4w  {7.0: 66534, 5.0: 64588, 8.0: 63581}

missing_count missing_percentage mean    std \
date_of_birth_distinct_emails_4w      0          0.0  9.5  5.04

min median    max corr_with_target
date_of_birth_distinct_emails_4w  0.0    9.0  39.0          -0.04

```



Feature: 'credit_risk_score'

Description: Internal score of application risk. Ranges between [-191, 389].

```

dtype    count  unique  \
credit_risk_score  int64  800000    550

```

```

top_value_counts  missing_count  \
credit_risk_score  {116.0: 5436, 115.0: 5417, 110.0: 5402}    0

```

```


missing_percentage  mean    std    min  median    max  \
credit_risk_score      0.0  131.01  69.68  -170.0  122.0  389.0

```

```

corr_with_target
credit_risk_score      0.07

```



Feature: 'bank_months_count'

Description: How old is previous account (if held) in months. Ranges between [-1, 32] months (-1 is a missing value).

```

dtype    count  unique  \

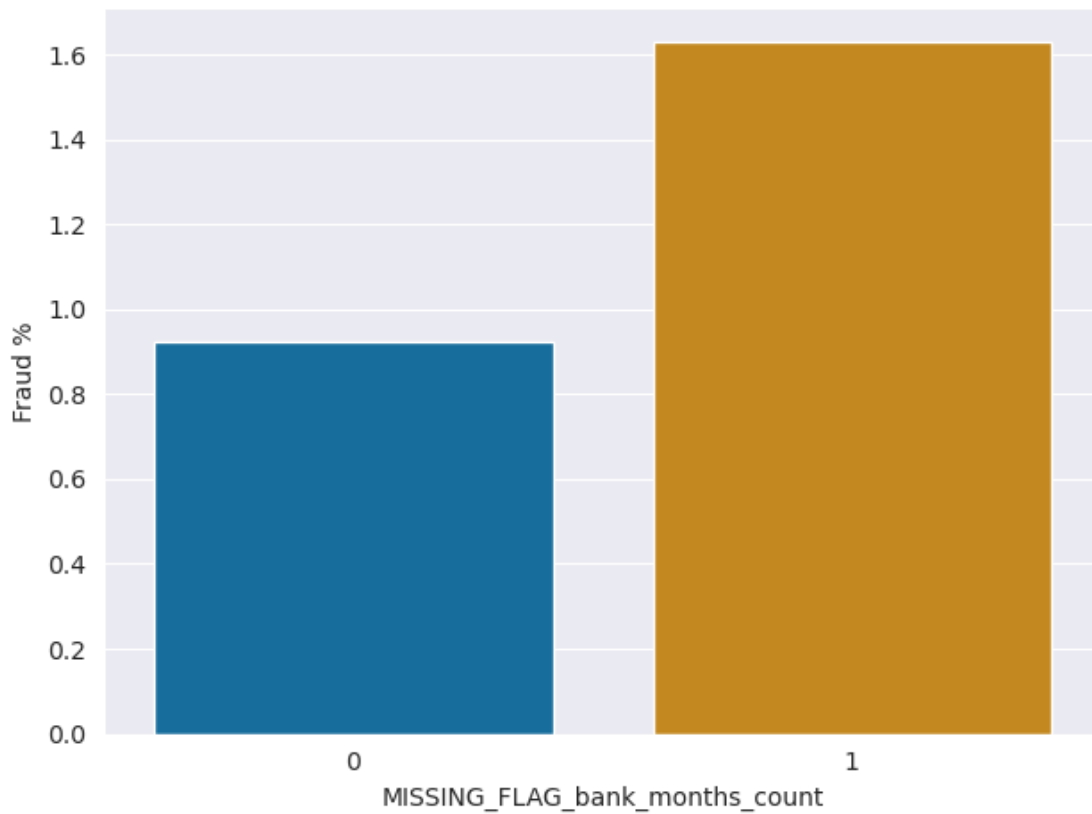

```

bank_months_count float64 597396 32

top_value_counts missing_count \
bank_months_count {1.0: 155878, 28.0: 64172, 15.0: 47326} 202604

missing_percentage mean std min median max \
bank_months_count 25.33 14.87 11.53 1.0 15.0 32.0

corr_with_target
bank_months_count 0.02



Feature: 'session_length_in_minutes'


Description: Length of user session in banking website in minutes.
Ranges between [-1, 107] minutes (-1 is a missing value).

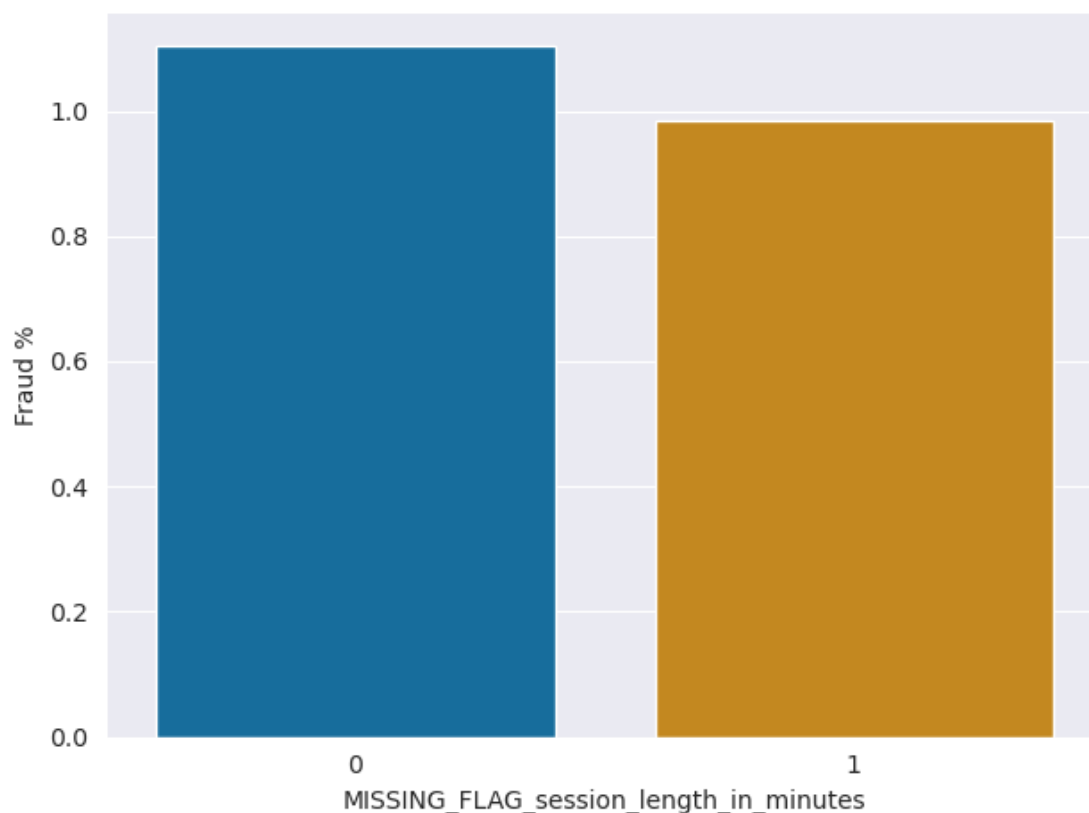
	dtype	count	unique	\
session_length_in_minutes	float64	798377	796390	

	top_value_counts	missing_count	\
session_length_in_minutes	{4.74: 3, 4.19: 3, 4.64: 3}	1623	

	missing_percentage	mean	std	min	median	max	\
session_length_in_minutes	0.2	7.57	8.04	0.0	5.12	85.9	

	corr_with_target
session_length_in_minutes	0.01





1.2 Discrete / Binary features

Applying summary function to all discrete / binary features:


```
[ ]: for feat in disc_feats:
    print(f"\033[1m Feature:\033[0m '{feat}'\n")
    print(f"\033[1m Description:\033[0m {feature_descriptions[feat]}\n")
    summary(feat)
    print('-'*45, '\n')
```

Feature: 'income'

Description: Annual income of the applicant (in decile form). Ranges between [0.1, 0.9].

	dtype	count	unique	missing_count	missing_percentage	mean	std	\
income	float64	800000	9	0	0.0	0.56	0.29	

	min	median	max	cv	corr_with_target
income	0.1	0.6	0.9	0.52	0.05




Feature: 'customer_age'

Description: Applicant's age in years, rounded to the decade. Ranges between [10, 90] years.

	dtype	count	unique	missing_count	missing_percentage	mean	\
customer_age	int64	800000	9	0	0.0	33.7	

	std	min	median	max	cv	corr_with_target
customer_age	12.03	10.0	30.0	90.0	0.36	0.06




Feature: 'email_is_free'

Description: Domain of application email (either free or paid).

	dtype	count	unique	missing_count	missing_percentage	mean	\
email_is_free	int64	800000	2	0	0.0	0.53	

	std	min	median	max	cv	corr_with_target
email_is_free	0.5	0.0	1.0	1.0	0.94	0.03




Feature: 'phone_home_valid'

Description: Validity of provided home phone.

	dtype	count	unique	missing_count	missing_percentage	mean	\
phone_home_valid	int64	800000	2	0	0.0	0.42	

	std	min	median	max	cv	corr_with_target
phone_home_valid	0.49	0.0	0.0	1.0	1.18	-0.03

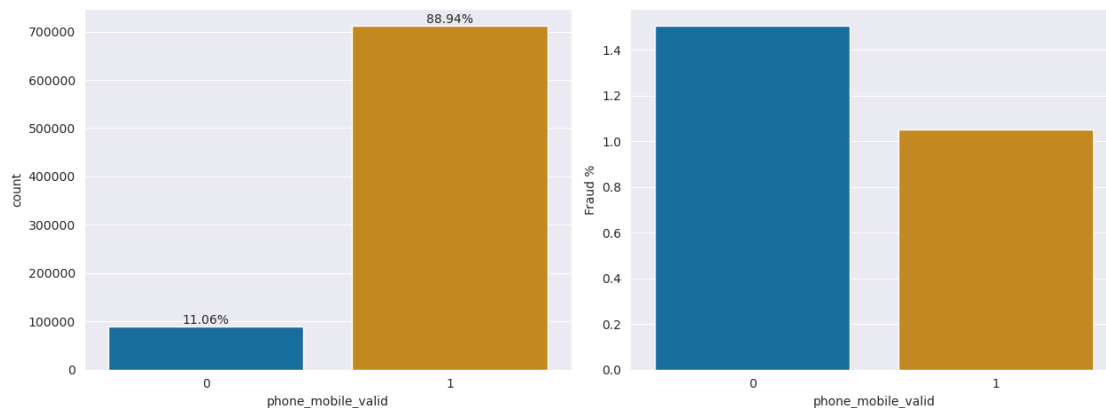


Feature: 'phone_mobile_valid'

Description: Validity of provided mobile phone.

	dtype	count	unique	missing_count	missing_percentage	\
phone_mobile_valid	int64	800000	2	0	0.0	

	mean	std	min	median	max	cv	corr_with_target
phone_mobile_valid	0.89	0.31	0.0	1.0	1.0	0.35	-0.01




Feature: 'has_other_cards'

Description: If applicant has other cards from the same banking company.

	dtype	count	unique	missing_count	missing_percentage	mean	\
has_other_cards	int64	800000	2	0	0.0	0.22	


	std	min	median	max	cv	corr_with_target
has_other_cards	0.42	0.0	0.0	1.0	1.87	-0.04



Feature: 'proposed_credit_limit'

Description: Applicant's proposed credit limit. Ranges between [200, 2000].

	dtype	count	unique	missing_count	\	
proposed_credit_limit	float64	800000	12	0		
	missing_percentage	mean	std	min	median	\
proposed_credit_limit	0.0	516.09	487.74	190.0	200.0	
	max	cv	corr_with_target			
proposed_credit_limit	2100.0	0.95	0.07			




Feature: 'foreign_request'

Description: If origin country of request is different from bank's country.

	dtype	count	unique	missing_count	missing_percentage	mean	\
foreign_request	int64	800000	2	0	0.0	0.03	

	std	min	median	max	cv	corr_with_target
foreign_request	0.16	0.0	0.0	1.0	6.21	0.02




Feature: 'keep_alive_session'

Description: User option on session logout.

	dtype	count	unique	missing_count	missing_percentage	\
keep_alive_session	int64	800000	2	0	0.0	


	mean	std	min	median	max	cv	corr_with_target
keep_alive_session	0.58	0.49	0.0	1.0	1.0	0.86	-0.05

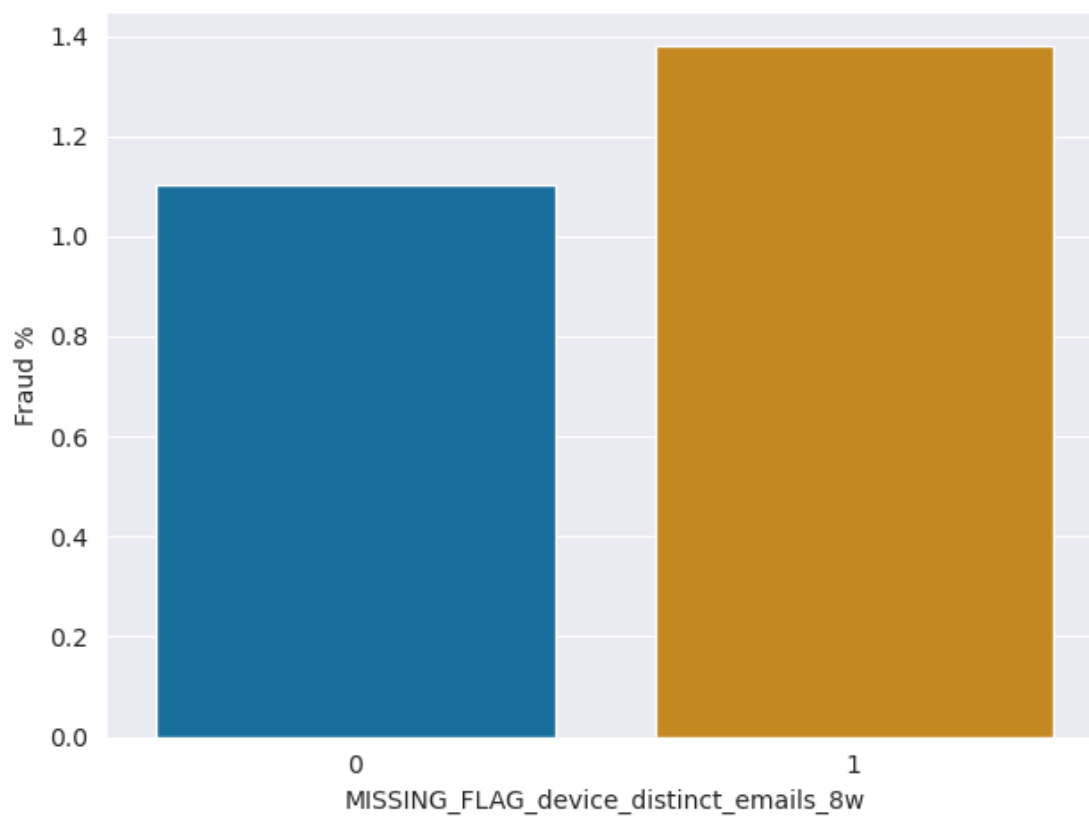


Feature: 'device_distinct_emails_8w'

Description: Number of distinct emails in banking website from the used device in last 8 weeks. Ranges between [-1, 2] emails (-1 is a missing value).

	dtype	count	unique	missing_count	\			
device_distinct_emails_8w	float64	799710	3	290				
	missing_percentage	mean	std	min	median	max	\	
device_distinct_emails_8w	0.04	1.02	0.18	0.0	1.0	2.0		
	cv	corr_with_target						
device_distinct_emails_8w	0.17	0.04						






Feature: 'device_fraud_count'

Description: Number of fraudulent applications with used device. Ranges between [0, 1].

	dtype	count	unique	missing_count	missing_percentage	\
device_fraud_count	int64	800000	1	0	0.0	


	mean	std	min	median	max	cv	corr_with_target
device_fraud_count	0.0	0.0	0.0	0.0	0.0	nan	nan



Feature: 'month'

Description: Month where the application was made. Ranges between [0, 7].

	dtype	count	unique	missing_count	missing_percentage	mean	std	min	\
month	int64	800000	8	0	0.0	3.29	2.21	0.0	
	median	max	cv	corr_with_target					
month	3.0	7.0	0.67	0.01					



1.3 Categorical features


Applying summary function to all categorical features:


```
[ ]: for feat in cat_feats:
    print(f"\033[1m Feature:\033[0m '{feat}'\n")
    print(f'\033[1m Description:\033[0m {feature_descriptions[feat]}')
    summary(feat)
    print('-'*45, '\n')
```

Feature: 'payment_type'

Description: Credit payment plan type. 5 possible (anonymized) values.


	dtype	count	unique	missing_count	missing_percentage
payment_type	object	800000	5	0	0.0



Feature: 'employment_status'

Description: Employment status of the applicant. 7 possible (anonymized) values.


	dtype	count	unique	missing_count	missing_percentage
employment_status	object	800000	7	0	0.0



Feature: 'housing_status'

Description: Current residential status for applicant. 7 possible (anonymized) values.


	dtype	count	unique	missing_count	missing_percentage
housing_status	object	800000	7	0	0.0



Feature: 'source'

Description: Online source of application. Either browser (INTERNET) or app (TELEAPP).


	dtype	count	unique	missing_count	missing_percentage
source	object	800000	2	0	0.0



Feature: 'device_os'

Description: Operative system of device that made request. Possible values are: Windows, macOS, Linux, X11, or other.

	dtype	count	unique	missing_count	missing_percentage
device_os	object	800000	5	0	0.0



1.4 Correlation coefficients

Correlation between each feature and the target (fraud or not):

```
[ ]: # Correlation coefficient between features and target

fig = plt.figure(figsize=(9.6,7.2))

corr_target = pd.DataFrame(pd.get_dummies(X_train.
    ↳drop('device_fraud_count',axis=1))
                           .corrwith(y_train)).reset_index()

corr_target.columns = ['feature','target_correlation']
corr_target = corr_target.sort_values(by='target_correlation')

ax = sns.
    ↳barplot(data=corr_target,x='feature',y='target_correlation',palette=palette)
ax.set_title('Feature correlations with target')
```

```

ax.set_xticklabels(ax.get_xticklabels(), rotation=90, ha="right")
ax.set_ylabel('Correlation with target')

fig.tight_layout()
plt.show()

```

